Reflection on AI-Sycophancy and AI-Psychosis

AI-Sycophancy

Working on this homework gave me a direct experience of AI sycophancy. When I asked the AI agent to evaluate companies, it sometimes told me what seemed "reasonable" rather than what the data actually showed. For example, when a company's management made vague positive statements about product quality, the agent initially tended to agree with them — giving PASS — even when the return rate data contradicted the claim. The AI was "agreeing" with the document rather than critically cross-checking it.

This mirrors a real danger in investment analysis: an AI that validates management narratives instead of verifying them is worse than useless. It creates false confidence. The fix was to force the agent to compute specific numerical thresholds before accepting any qualitative claim as true.

AI-Psychosis

I also experienced AI-psychosis when the agent became too aggressive. After I instructed it to "actively hunt for false statements," it started finding problems everywhere — flagging normal 2-3% return rates as fraud, treating routine R&D fluctuations as evidence of misconduct, and misinterpreting revenue recognition differences as lies. The agent lost touch with what "normal" looked like and started hallucinating red flags that didn't exist.

This is psychosis in a real sense: the model's internal reasoning diverged from external reality. It was generating a coherent-sounding story that was factually wrong.

The Connection

Sycophancy and psychosis are opposites but both stem from the same root — the model optimizing for something other than truth. Sycophancy optimizes for agreement; psychosis optimizes for finding patterns even where none exist. Both produce unreliable outputs. The lesson: AI agents need grounded rules and numerical anchors, not just language-level instructions.